

regression_model

August 7, 2024

0.1 Data processing

0.1.1 There are three techniques to solve the missing values' problem in order to find out the most accurate features, and they are:

0.1.2 * Dropping

0.1.3 * Numerical imputation

0.1.4 * Categorical imputation

```
[1]: import pandas as pd
import numpy as np
from scipy import stats

[2]: df = pd.read_csv('/home/student/220962344_ml_lab/Week3/diabetes_csv.csv')
df

## To check for missing values
# print(df.isnull())
df['Has_Missing'] = df.isnull().any(axis=1)

if df['Has_Missing'].any() == True:
    print(yes)
# This code takes the mean of the columns and retains the cols with more than
↳ 60% non-missing value
# Use axis = 1 in mean for rows
threshold = 60
df = df.loc[df.isnull().mean(axis=1) < threshold] # rows
df = df[df.columns[df.isnull().mean() < threshold]] # cols
df

## We use imputation to prevent low training size
# Replace the missing data with a relevant value
df = df.fillna(0)
df = df.fillna(df.median())

## Outlier Identification
# * Z-Score and IQR are good for identifying outliers in univariate data.
# * Box Plots and Scatter Plots offer a visual approach to detecting outliers.
```

```

# * Machine Learning Methods: Isolation Forest and Local Outlier Factor are
  ↪useful for more complex datasets.
df['Z-Score_glucose'] = stats.zscore(df['Glucose'])
df_filtered = df[abs(df['Z-Score']) <= 3]
print(df_filtered.head())

Q1 = df['Glucose'].quantile(0.25)
Q3 = df['Glucose'].quantile(0.75)
IQR = Q3 - Q1

df_filtered = df[(df['Glucose'] >= Q1) & (df['Glucose'] <= Q3)]

## Overfitting
# We can use binning to remove rows ith less importance
df['bin'] = pd.cut(df['Glucose'], bins=[100,250,400,500], labels=["Lowest",
  ↪"Mid", "High"])

```

```

-----
KeyError                                Traceback (most recent call last)
File /usr/lib/python3/dist-packages/pandas/core/indexes/base.py:3791, in Index.
  ↪get_loc(self, key)
    3790 try:
-> 3791     return self._engine.get_loc(casted_key)
    3792 except KeyError as err:

File /usr/lib/python3/dist-packages/pandas/_libs/index.pyx:152, in pandas._libs
  ↪index.IndexEngine.get_loc()

File /usr/lib/python3/dist-packages/pandas/_libs/index.pyx:181, in pandas._libs
  ↪index.IndexEngine.get_loc()

File pandas/_libs/hashtable_class_helper.pxi:7080, in pandas._libs.hashtable.
  ↪PyObjectHashTable.get_item()

File pandas/_libs/hashtable_class_helper.pxi:7088, in pandas._libs.hashtable.
  ↪PyObjectHashTable.get_item()

KeyError: 'Z-Score'

```

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call last)
Cell In[2], line 27
    22 ## Outlier Identification
    23 # * Z-Score and IQR are good for identifying outliers in univariate dat.

```

```

24 # * Box Plots and Scatter Plots offer a visual approach to detecting
↳ outliers.
25 # * Machine Learning Methods: Isolation Forest and Local Outlier Factor
↳ are useful for more complex datasets.
26 df['Z-Score_glucose'] = stats.zscore(df['Glucose'])
--> 27 df_filtered = df[abs(df['Z-Score']) <= 3]
28 print(df_filtered.head())
30 Q1 = df['Glucose'].quantile(0.25)

```

File /usr/lib/python3/dist-packages/pandas/core/frame.py:3893, in DataFrame.

```

↳ __getitem__(self, key)
3891 if self.columns.nlevels > 1:
3892     return self._getitem_multilevel(key)
-> 3893 indexer = self.columns.get_loc(key)
3894 if is_integer(indexer):
3895     indexer = [indexer]

```

File /usr/lib/python3/dist-packages/pandas/core/indexes/base.py:3798, in Index.

```

↳ get_loc(self, key)
3793 if isinstance(casted_key, slice) or (
3794     isinstance(casted_key, abc.Iterable)
3795     and any(isinstance(x, slice) for x in casted_key)
3796 ):
3797     raise InvalidIndexError(key)
-> 3798 raise KeyError(key) from err
3799 except TypeError:
3800     # If we have a listlike key, _check_indexing_error will raise
3801     # InvalidIndexError. Otherwise we fall through and re-raise
3802     # the TypeError.
3803     self._check_indexing_error(key)

```

KeyError: 'Z-Score'

```

[ ]: ## We do encoding for non numeric data
from sklearn.preprocessing import LabelEncoder
# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Fit and transform the data
df['Priority_Encoded'] = label_encoder.fit_transform(df['__name__'])

```

0.2 Lab Questions

```

[3]: '''
Consider the hepatitis/ pima-indians-diabetes csv file, perform the following
↳ date pre-processing.

```

1. Load data in Pandas.
2. Drop columns that aren't useful.
3. Drop rows with missing values.
4. Create dummy variables.
5. Take care of missing data.
6. Convert the data frame to NumPy.
7. Divide the data set into training data and test data.

```
'''
import pandas as pd
import numpy as np

df = pd.read_csv('hepatitis_csv.csv')

# the class is the target variable, live or die
df.head()

df_cleaned = df.dropna()
df_cleaned.head()

df_dummies = pd.get_dummies(df_cleaned, columns=['sex', 'class'])
df_dummies.head()

numpy = df_dummies.to_numpy()
print(numpy)

# shuffle
df_dummies = df_dummies.sample(frac=1)

# Define the split ratio
train_ratio = 0.8
train_size = int(df_dummies.shape[0] * train_ratio)

# Split into training and testing sets
train_df = df_dummies[:train_size]
test_df = df_dummies[train_size:]

print("Train")
print(train_df)

print("Test:")
print(test_df)
```

```
[[34 True False ... False False True]
 [39 False True ... False False True]
 [32 True True ... False False True]
 ...
 [31 False False ... False False True]
```

```
[53 False False ... True False True]
[43 True False ... False True False]]
```

Train

	age	steroid	antivirals	fatigue	malaise	anorexia	liver_big	liver_firm	\
96	30	False	False	True	True	False	True	True	
23	42	True	False	False	False	False	True	False	
125	34	True	False	True	True	True	False	True	
25	27	False	False	True	True	False	True	False	
109	33	False	False	True	True	False	True	False	
..	
133	72	True	True	True	False	False	True	True	
60	37	True	False	False	False	False	True	False	
54	30	True	False	True	False	False	True	False	
77	34	False	True	False	False	False	True	True	
135	25	True	False	True	False	False	False	True	

	spleen_palpable	spiders	...	bilirubin	alk_phosphate	sgot	albumin	\
96	False	True	...	0.8	147.0	128.0	3.9	
23	False	False	...	0.9	60.0	63.0	4.7	
125	False	True	...	0.7	70.0	24.0	4.1	
25	False	False	...	0.8	95.0	46.0	3.8	
109	False	False	...	0.7	63.0	80.0	3.0	
..	
133	False	False	...	1.0	115.0	52.0	3.4	
60	False	False	...	0.7	26.0	58.0	4.5	
54	False	False	...	0.7	50.0	78.0	4.2	
77	False	False	...	0.6	30.0	24.0	4.0	
135	True	True	...	1.3	181.0	181.0	4.5	

	protime	histology	sex_female	sex_male	class_die	class_live
96	100.0	True	True	False	False	True
23	47.0	False	True	False	False	True
125	100.0	True	False	True	False	True
25	100.0	False	True	False	False	True
109	31.0	True	True	False	True	False
..
133	50.0	True	True	False	False	True
60	100.0	False	True	False	False	True
54	74.0	False	True	False	False	True
77	76.0	False	False	True	False	True
135	57.0	True	True	False	False	True

[64 rows x 22 columns]

Test:

	age	steroid	antivirals	fatigue	malaise	anorexia	liver_big	liver_firm	\
89	38	False	False	True	True	True	False	True	
98	47	True	False	False	False	False	True	False	
138	47	True	False	True	True	False	True	True	

19	38	False	True	False	False	False	False	True
75	32	False	True	True	True	False	True	False
38	42	False	False	False	False	False	True	False
53	40	True	True	True	False	False	True	True
94	59	False	False	True	True	False	True	True
39	65	True	False	True	True	False	True	True
15	38	False	False	True	True	True	True	False
85	28	False	False	True	True	True	True	True
5	34	True	False	False	False	False	True	False
124	50	True	False	False	False	False	True	True
33	26	False	False	False	False	False	True	True
43	56	False	False	True	False	False	True	False
17	40	False	False	True	False	False	True	True

	spleen_palpable	spiders	...	bilirubin	alk_phosphate	sgot	albumin	\
89	False	False	...	0.6	76.0	18.0	4.4	
98	False	True	...	2.0	84.0	23.0	4.2	
138	False	True	...	1.0	166.0	30.0	2.6	
19	False	False	...	0.7	70.0	28.0	4.2	
75	False	False	...	1.0	55.0	45.0	4.1	
38	False	False	...	1.0	85.0	14.0	4.0	
53	True	False	...	1.2	85.0	31.0	4.0	
94	True	True	...	1.5	107.0	157.0	3.6	
39	True	True	...	0.3	180.0	53.0	2.9	
15	False	False	...	2.0	72.0	89.0	2.9	
85	False	False	...	1.6	44.0	123.0	4.0	
5	False	False	...	0.9	95.0	28.0	4.0	
124	True	True	...	1.0	85.0	75.0	4.0	
33	False	False	...	0.5	135.0	29.0	3.8	
43	False	False	...	0.7	71.0	18.0	4.4	
17	False	False	...	0.6	62.0	166.0	4.0	

	protime	histology	sex_female	sex_male	class_die	class_live
89	84.0	True	True	False	False	True
98	66.0	True	True	False	True	False
138	31.0	True	True	False	True	False
19	62.0	False	True	False	False	True
75	56.0	False	True	False	False	True
38	100.0	False	True	False	False	True
53	100.0	False	True	False	False	True
94	38.0	True	True	False	True	False
39	74.0	True	True	False	False	True
15	46.0	False	True	False	False	True
85	46.0	False	True	False	False	True
5	75.0	False	True	False	False	True
124	72.0	True	True	False	False	True
33	60.0	False	False	True	False	True
43	100.0	False	True	False	False	True

17 63.0 False True False False True

[16 rows x 22 columns]

```
[4]: '''
2. a. Construct a CSV file with the following attributes:
Study time in hours of ML lab course (x)
Score out of 10 (y)
The dataset should contain 10 rows.
b. Create a regression model and display the following:
Coefficients: B0 (intercept) and B1 (slope)
RMSE (Root Mean Square Error)
Predicted responses
c. Create a scatter plot of the data points in red color and plot the graph of
    ↪ x vs. predicted y in blue color.
d. Implement the model using two methods:
Pedhazur formula (intuitive)
Calculus method (partial derivatives, refer to class notes)
e. Compare the coefficients obtained using both methods and compare them with
    ↪ the analytical solution.
f. Test your model to predict the score obtained when the study time of a
    ↪ student is 10 hours.
Note: Do not use scikit-learn.
'''

import numpy as np
import pandas as pd

study_time = np.random.randint(1,6,100)

score = np.random.randint(0,11,100)

dic = {
    "study_time" : study_time,
    "Score" : score
}
df = pd.DataFrame(dic)
df.head()
```

```
[4]:   study_time  Score
0          3       7
1          4       3
2          2       0
3          4       1
4          1       5
```

```
[6]: import numpy as np

class LinearRegression:
    def __init__(self):
        self.weight = None
        self.bias = None
        self.epochs = []
        self.loss = []
    def fit(self, X, y, epochs=1000, alpha=0.01):
        """
        Train the linear regression model using gradient descent.

        Parameters:
        X (numpy array): Input feature of shape (num_samples,)
        y (numpy array): Target values of shape (num_samples,)
        epochs (int): Number of iterations for gradient descent
        alpha (float): Learning rate
        """
        num_samples = len(X)
        # Initialize weight and bias
        self.weight = 0.0
        self.bias = 0.0

        # Gradient descent
        for epoch in range(epochs):
            # Compute predictions
            y_pred = self.predict(X)
            # Compute gradients
            weight_gradient = -2 * np.sum((y - y_pred) * X) / num_samples
            bias_gradient = -2 * np.sum(y - y_pred) / num_samples
            # Update parameters
            self.weight -= alpha * weight_gradient
            self.bias -= alpha * bias_gradient

            # Optionally print loss for every 100 epochs
            if epoch % 100 == 0:
                loss = np.mean((y_pred - y) ** 2)
                print(f"Epoch {epoch}, Loss: {loss}")
                self.epochs.append(epoch)
                self.loss.append(loss)

    def predict(self, X):
        """
        Predict the target values for the input feature.

        Parameters:
```



```

X (numpy array): Input feature of shape (num_samples,)

Returns:
numpy array: Predicted target values
'''
if self.weight is None or self.bias is None:
    raise ValueError("Model has not been trained yet.")
return self.weight * X + self.bias

def mean_squared_error(self, y_true, y_pred):
    '''
    Calculate the mean squared error between true and predicted values.

    Parameters:
    y_true (numpy array): True target values
    y_pred (numpy array): Predicted target values

    Returns:
    float: Mean squared error
    '''
    return np.mean((y_true - y_pred) ** 2)

```

```

[7]: import numpy as np

# Generate synthetic data
np.random.seed(42)
X = np.random.rand(100) # 100 samples of input feature
true_weight = 3.5
true_bias = -2.0
y = np.random.rand(100)

# Create and train the model
model = LinearRegression()
model.fit(X, y, epochs=1000, alpha=0.01)

# Predict using the trained model
y_pred = model.predict(X)

# Calculate and print the mean squared error
mse = model.mean_squared_error(y, y_pred)
print(f"Mean Squared Error: {mse}")

# Print learned parameters
print(f"Learned weight: {model.weight}")
print(f"Learned bias: {model.bias}")

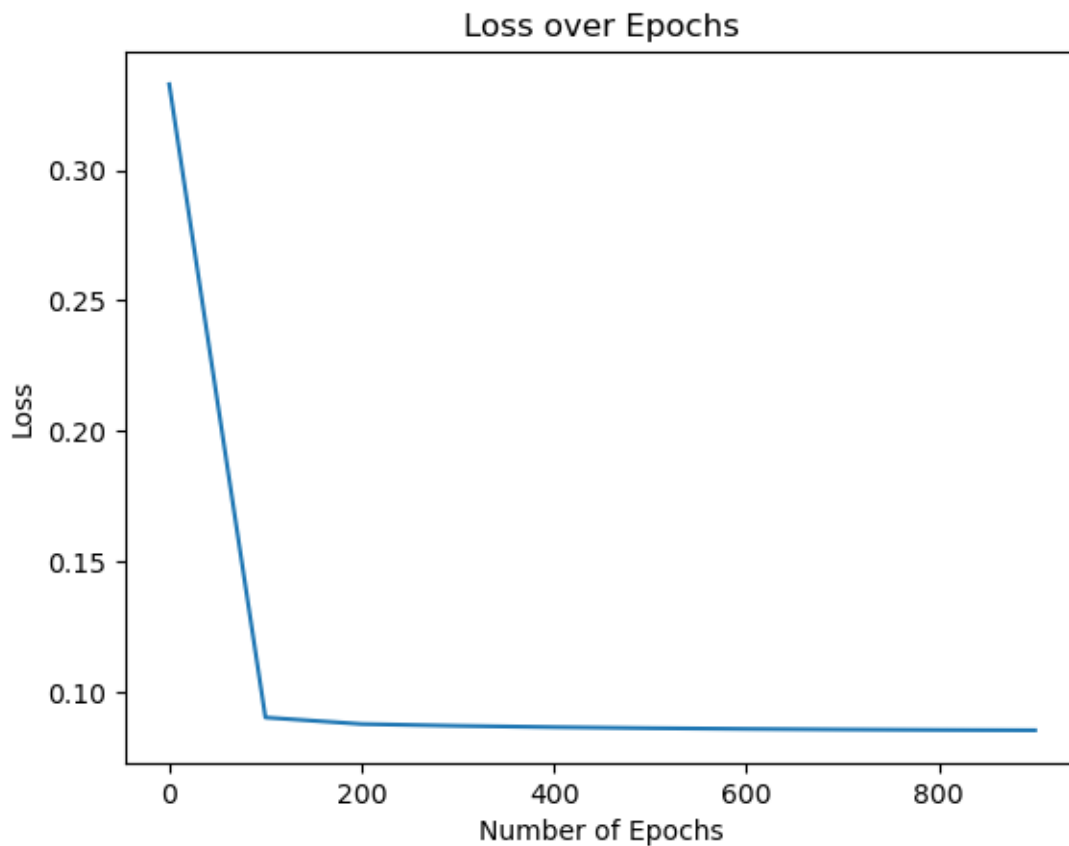
```

Epoch 0, Loss: 0.33289148886695125
Epoch 100, Loss: 0.09022597115000433

Epoch 200, Loss: 0.08771998112325295
Epoch 300, Loss: 0.08703019843125777
Epoch 400, Loss: 0.08651855289171552
Epoch 500, Loss: 0.08613318908361896
Epoch 600, Loss: 0.08584289931278948
Epoch 700, Loss: 0.08562422737112567
Epoch 800, Loss: 0.08545950431374577
Epoch 900, Loss: 0.08533542033009127
Mean Squared Error: 0.0852419492930398
Learned weight: 0.023128772113414613
Learned bias: 0.4849278695821864

```
[8]: import matplotlib.pyplot as plt # Correct import statement

# Assuming `model` has `epochs` and `loss` attributes
plt.plot(model.epochs, model.loss)
plt.xlabel('Number of Epochs')
plt.ylabel('Loss')
plt.title('Loss over Epochs')
plt.show()
```



[]:

[]: